

NYU HPC Cloud Bursting Tutorial

Course: CSCI-UA.0480-051 (Fall 2025)

Contents

1 Your Resource Allocation	3
1.1 Available Partitions	3
2 Prerequisites	3
2.1 Network Connection	3
3 Quick Start Guide	3
3.1 Method 1: Using Open OnDemand (Recommended)	3
3.1.1 Step 1: Access Open OnDemand	3
3.1.2 Step 2: Upload Your Files	3
3.1.3 Step 3: Open a Terminal	4
3.2 Method 2: Using Terminal/SSH	4
3.2.1 SSH Connection	4
3.3 Method 3: Using VS Code Remote SSH	4
3.3.1 Step 1: Install VS Code Extension	4
3.3.2 Step 2: Configure SSH	4
3.3.3 Step 3: Connect via VS Code	5
4 Running Interactive Jobs	5
4.1 CPU Job for MPI Programs	5
4.2 GPU Jobs for CUDA Programs	6
5 Setting Up Singularity for MPI Programs	7
5.1 One-Time Setup	7
5.1.1 Step 1: Export PATH and Navigate to Scratch Directory	7
5.1.2 Step 2: Download Singularity Image	7
5.1.3 Step 3: Download and Extract Overlay File	8
5.1.4 Step 4: Launch Singularity Container	8
5.1.5 Step 5: Install Miniconda (One-Time Only)	8
5.1.6 Step 6: Install OpenMPI via Conda (One-Time Only)	8
5.2 Regular Usage After Setup	8
6 Running MPI Programs	9
6.1 Complete MPI Workflow	9
6.2 Important Notes for MPI	9
7 Running CUDA Programs	10
7.1 Basic CUDA Workflow	10
7.2 Available CUDA Modules	10

8 Example Programs	10
8.1 Example 1: Basic MPI Hello World (WITH Singularity)	10
8.2 Example 2: CUDA Vector Addition	11
8.3 Example 4: Complete Singularity Setup from Scratch	13
9 Batch Job Scripts	14
9.1 MPI Batch Job	14
9.2 CUDA Batch Job	14
10 Troubleshooting	14
10.1 Common MPI Errors and Solutions	14
10.1.1 Error: libpmix.so.2: cannot open shared object file	14
10.1.2 FATAL: while loading overlay images: failed to open overlay image	15
10.1.3 Error: Not enough slots available	15
10.1.4 Error: File not found when copying Singularity image	15
10.1.5 Error: command not found: conda	15
10.2 For CUDA Programs	15
10.2.1 Error: nvcc: command not found	15
10.2.2 Error: No CUDA-capable device	16
10.3 SSH Host Key Issues	16
11 Debugging Tips	16
11.1 For MPI Programs	16
11.2 For CUDA Programs	16
11.3 General Debugging	17
12 Time and Resource Management	17
12.1 GPU Hour Budget	17
12.2 Best Practices	17
13 Quick Reference	17
13.1 Essential Commands	17
13.2 MPI Workflow Summary	18
13.3 CUDA Workflow Summary	18
14 Summary	18
14.1 Key Takeaways	18
14.2 Recommended Workflow	19
14.3 Important Links	19

1 Your Resource Allocation

Each student has been assigned:

- **Slurm Account:** csci_ua_0480_051-2025fa
- **GPU Hours:** 300 hours (18,000 minutes)
- **CPU Time:** Sufficient for coursework

1.1 Available Partitions

- **interactive** - For CPU-only interactive jobs
- **n2c48m24** - CPU partition
- **g2-standard-12, g2-standard-24, g2-standard-48** - GPU partitions with L4 GPUs
- **c12m85-a100-1, c24m170-a100-2** - A100 GPU partitions
- **n1s8-t4-1** - T4 GPU partition

2 Prerequisites

2.1 Network Connection

VPN Requirement

You need to connect to NYU VPN **only if you are NOT on NYU campus WiFi**. If you are on campus, you can skip the VPN setup.

For off-campus access: Download NYU VPN from <https://www.nyu.edu/life/information-technology/infrastructure/network-services/vpn.html>

3 Quick Start Guide

3.1 Method 1: Using Open OnDemand (Recommended)

This is the easiest way to get started and manage files - no complex command line setup required!

3.1.1 Step 1: Access Open OnDemand

- Connect to NYU VPN (if off-campus) or NYU WiFi (if on campus)
- Go to: <https://ood-burst-001.hpc.nyu.edu/>
- Log in with your NYU credentials (NetID and password)

What You Can Do with Open OnDemand

From the OOD server, you can:

- Launch compute nodes without logging into Greene cluster
- Run Jupyter notebooks (only if needed)
- Open terminal sessions directly in your browser
- Transfer files between your local computer and HPC
- Manage files with a graphical interface

3.1.2 Step 2: Upload Your Files

1. From the Open OnDemand homepage, click **Files** → **Home Directory**
2. Click the **Upload** button

3. Select your MPI or CUDA source files from your computer
4. Your files will be uploaded to your home directory on HPC

3.1.3 Step 3: Open a Terminal

1. Click **Clusters → Greene Shell Access**
2. This opens a terminal in your browser
3. Run **ssh burst** to connect to Cloud Burst
4. Now you can compile and run your programs

3.2 Method 2: Using Terminal/SSH

For users comfortable with command line access.

3.2.1 SSH Connection

Step 1: SSH to Greene Jump Host

```
1 ssh netID@gw.hpc.nyu.edu
```

Replace **netID** with your actual NYU NetID. For example, if your NetID is **ab1234**:

```
1 ssh ab1234@gw.hpc.nyu.edu
```

Step 2: SSH to Greene Cluster

```
1 ssh netID@greene.hpc.nyu.edu
```

Step 3: SSH to Cloud Burst

```
1 ssh burst
```

Note

When connecting for the first time, you may see: “This key is not known by any other names. Are you sure you want to continue connecting?”

- Type **yes** and press Enter

3.3 Method 3: Using VS Code Remote SSH

VS Code provides an excellent development environment with syntax highlighting, IntelliSense, and integrated terminal.

3.3.1 Step 1: Install VS Code Extension

1. Download VS Code from <https://code.visualstudio.com/>
2. Install the “Remote - SSH” extension

3.3.2 Step 2: Configure SSH

Click on the icon on the bottom left corner: follow fig 1 then fig 2

Then add this to your SSH config file (`~/.ssh/config`):

```
1 # NYU HPC Gateway
2 Host greene.hpc.nyu.edu
3   HostName greene.hpc.nyu.edu
4   User <net_id>
```

```

5 Host greene.hpc.nyu.edu dtn.hpc.nyu.edu gw.hpc.nyu.edu
6   StrictHostKeyChecking no
7   ServerAliveInterval 60
8   ForwardAgent yes
9   StrictHostKeyChecking no
10  UserKnownHostsFile /dev/null
11  LogLevel ERROR
12
13
14 Host hpcgwtunnel
15   HostName gw.hpc.nyu.edu
16   ForwardX11 no
17   StrictHostKeyChecking no
18   LocalForward 8027 greene.hpc.nyu.edu:22
19   UserKnownHostsFile /dev/null
20   User <net_id>
21
22 Host greene
23   HostName localhost
24   Port 8027
25   ForwardX11 yes
26   StrictHostKeyChecking no
27   UserKnownHostsFile /dev/null
28   LogLevel ERROR
29   User <net_id>
30
31 Host burst.hpc.nyu.edu
32   HostName burst.hpc.nyu.edu
33   User <netid>
34   ServerAliveInterval 60
35   ForwardAgent yes

```

Replace <net_id> with your actual NetID.

3.3.3 Step 3: Connect via VS Code

1. Press F1 or Ctrl+Shift+P
2. Type “Remote-SSH: Connect to Host”
3. Select `greene.hpc.nyu.edu`
4. Enter your NYU password
5. Open folder: `/scratch/your_netid`
6. Start new terminal and run `ssh burst`. look at fig 3 for reference.
7. Edit, compile, and run code directly in VS Code

4 Running Interactive Jobs

4.1 CPU Job for MPI Programs

IMPORTANT: Specify CPUs

When running MPI programs, you **MUST** specify the number of CPUs you need using `--cpus-per-task=N` or `--ntasks=N`, where N is the number of MPI processes you plan to run. Otherwise, you'll only get 1 CPU and cannot run parallel programs.

For 8 MPI processes:

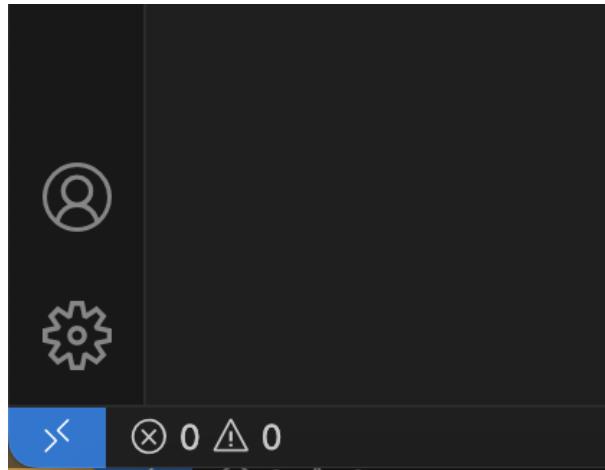


Figure 1: Click this icon to open a dialogue box

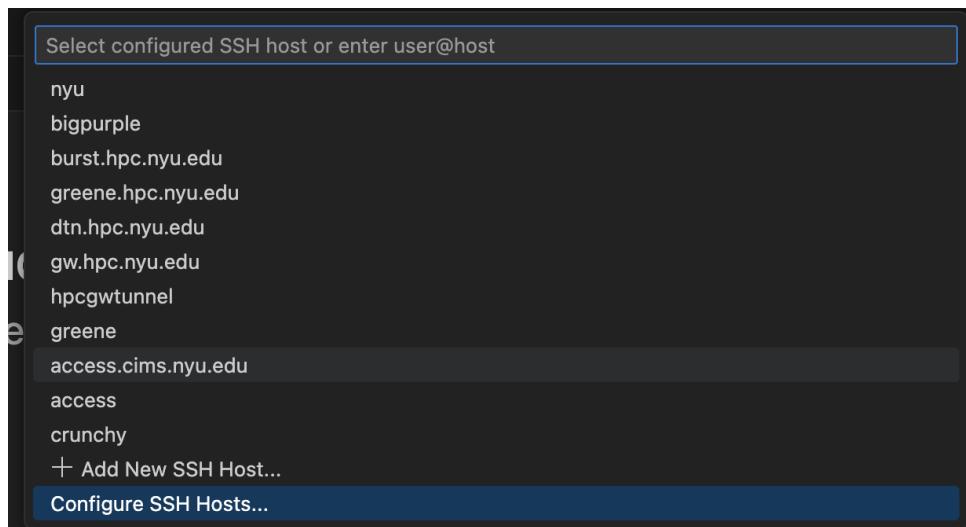


Figure 2: click on configure SSH Hosts (only for first time setup)

```
1 srun --account=csci_ua_0480_051-2025fa --partition=n2c48m24 --cpus-per-task=8 --time
   =04:00:00 --pty /bin/bash
```

Alternative with ntasks:

```
1 srun --account=csci_ua_0480_051-2025fa --partition=n2c48m24 --ntasks=2 --time=04:00:00
   --pty /bin/bash
```

4.2 GPU Jobs for CUDA Programs

1 L4 GPU for 4 Hours (Recommended):

```
1 srun --account=csci_ua_0480_051-2025fa --partition=g2-standard-12 --gres=gpu:1 --time
   =04:00:00 --pty /bin/bash
```

1 A100 GPU for 4 Hours:

```
1 srun --account=csci_ua_0480_051-2025fa --partition=c12m85-a100-1 --gres=gpu --time
   =04:00:00 --pty /bin/bash
```

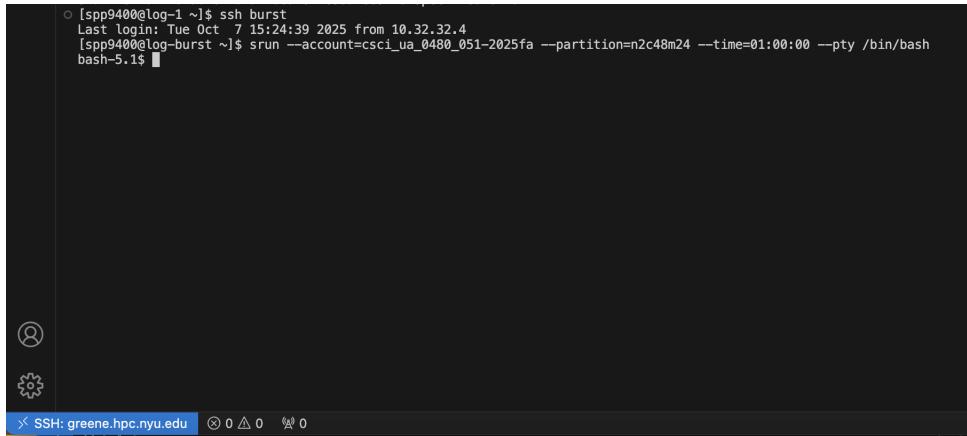


Figure 3: This is how it will look like

5 Setting Up Singularity for MPI Programs

REQUIRED for MPI Programs

Singularity is REQUIRED for MPI programs on burst nodes! The burst compute nodes have an incomplete MPI installation that causes `libpmix.so.2` errors. You must use Singularity containers to run MPI programs successfully.

5.1 One-Time Setup

You only need to do this setup once. After that, you can reuse the same Singularity container and overlay.

5.1.1 Step 1: Export PATH and Navigate to Scratch Directory

```

1 # Check if OpenMPI packages are available via the system package manager
2 which mpicc
3 ls /usr/lib64/openmpi/bin/
4 # If it exists in /usr/lib64/openmpi/bin/, you can add it to your PATH:
5 export PATH=/usr/lib64/openmpi/bin:$PATH
6 export LD_LIBRARY_PATH=/usr/lib64/openmpi/lib:$LD_LIBRARY_PATH
7
8 # Then try:
9 which mpicc
10 which mpirun
11 # After getting your compute node, go to scratch
12 cd /scratch/<your_netid> or cd ../../scratch/<your_netid>
13 # Example: cd /scratch/ab1234, based on your location: use pwd to check

```

5.1.2 Step 2: Download Singularity Image

```

1 # Download the Ubuntu Singularity image
2 scp -rp greene-dtn:/scratch/work/public/singularity/ubuntu-20.04.3.sif .

```

This will take a few minutes as the file is several GB.

5.1.3 Step 3: Download and Extract Overlay File

```
1 # Download overlay template
2 scp -rp greene-dtn:/scratch/work/public/overlay-fs-ext3/overlay-15GB-500K.ext3.gz .
3
4 # Extract it
5 gunzip overlay-15GB-500K.ext3.gz
```

This will take a few minutes as the file is several GB.

5.1.4 Step 4: Launch Singularity Container

```
1 # Start the container with overlay
2 singularity exec --overlay overlay-15GB-500K.ext3:rw ubuntu-20.04.3.sif /bin/bash
```

Your prompt will change to show you're inside Singularity (e.g., `Singularity>`).

5.1.5 Step 5: Install Miniconda (One-Time Only)

```
1 # Inside Singularity container
2 # Download Miniconda
3 wget https://repo.anaconda.com/miniconda/Miniconda3-latest-Linux-x86_64.sh
4
5 # Install to /ext3 (the overlay filesystem)
6 bash Miniconda3-latest-Linux-x86_64.sh -b -p /ext3/miniconda3
7
8 # Activate conda
9 source /ext3/miniconda3/bin/activate
```

5.1.6 Step 6: Install OpenMPI via Conda (One-Time Only)

```
1 # Still inside Singularity, with conda activated
2 # accepting the terms and conditions the first time
3 conda config --add channels https://repo.anaconda.com/pkgs/main
4 conda config --add channels https://repo.anaconda.com/pkgs/r
5 conda install gcc_linux-64 gxx_linux-64 openmpi -y
```

This installs a working MPI implementation that doesn't have the `libpmix` issues.

5.2 Regular Usage After Setup

After the one-time setup, your workflow for each session is:

```
1 # 1. Get compute node with enough CPUs
2 srun --account=csci_ua_0480_051-2025fa --partition=n2c48m24 --cpus-per-task=8 --time
   =04:00:00 --pty /bin/bash
3
4 # 2. Go to your scratch directory
5 cd /scratch/<your_netid>
6 # or
7 cd ../../scratch/<your_netid> if you are in home directory
8
9 # 3. Launch Singularity (note: using :ro for read-only)
10 singularity exec --overlay overlay-15GB-500K.ext3:ro ubuntu-20.04.3.sif /bin/bash
11
12 # 4. Activate conda
```

```

13 source /ext3/miniconda3/bin/activate
14
15 # 5. Now you can compile and run MPI programs!
16 mpicc -o hello_mpi hello_mpi.c
17 mpirun -np 2 ./hello_mpi

```

Pro Tip

Use `:ro` (read-only) instead of `:rw` when you don't need to install new software. This is safer and prevents accidental modifications.

6 Running MPI Programs

6.1 Complete MPI Workflow

```

1 # 1. SSH to burst
2 ssh <netid>@gw.hpc.nyu.edu
3 ssh <netid>@greene.hpc.nyu.edu
4 ssh burst
5
6 # 2. Request compute node with CPUs matching your MPI process count
7 srun --account=csci_ua_0480_051-2025fa --partition=n2c48m24 --cpus-per-task=8 --time
   =04:00:00 --pty /bin/bash
8
9 # 3. Go to scratch and launch Singularity
10 cd /scratch/<netid>
11 singularity exec --overlay overlay-15GB-500K.ext3:ro ubuntu-20.04.3.sif /bin/bash
12
13 # 4. Activate conda
14 source /ext3/miniconda3/bin/activate
15
16 # 5. Create your MPI program
17 vi hello_mpi.c
18
19 # 6. Compile with mpicc
20 mpicc -o hello_mpi hello_mpi.c
21
22 # 7. Run with mpirun
23 mpirun -np 2 ./hello_mpi
24
25 # 8. When done, exit
26 exit # Exit Singularity
27 exit # Exit compute node

```

6.2 Important Notes for MPI

- **Match CPUs to processes:** If you run `mpirun -np 8`, you need `--cpus-per-task=8` or more
- **Always use Singularity:** Running `mpicc/mpirun` directly on burst nodes will fail with `libpmix` errors
- **Activate conda:** Don't forget `source /ext3/miniconda3/bin/activate` inside Singularity
- **Work in /scratch:** Store your code in `/scratch/<netid>` for better performance

7 Running CUDA Programs

7.1 Basic CUDA Workflow

```
1 # 1. Request GPU node
2 srun --account=csci_ua_0480_051-2025fa --partition=g2-standard-12 --gres=gpu:1 --time
   =04:00:00 --pty /bin/bash
3
4 # 2. Load CUDA module
5 module load cuda/11.8.0
6
7 # 3. Verify GPU
8 nvidia-smi
9
10 # 4. Create your CUDA program
11 vi hello_cuda.cu
12
13 # 5. Compile
14 nvcc -o hello_cuda hello_cuda.cu
15
16 # 6. Run
17 ./hello_cuda
```

7.2 Available CUDA Modules

Check available CUDA versions:

```
1 module avail cuda
```

Common versions:

- cuda/11.8.0 - Recommended for most programs
- cuda/12.0 - For newer GPU features

8 Example Programs

8.1 Example 1: Basic MPI Hello World (WITH Singularity)

This is the complete, corrected workflow that actually works on burst nodes.

Step 1: Create the MPI program

```
1 // hello_mpi.c
2 #include <mpi.h>
3 #include <stdio.h>
4
5 int main(int argc, char** argv) {
6     MPI_Init(&argc, &argv);
7
8     int world_size, world_rank;
9     MPI_Comm_size(MPI_COMM_WORLD, &world_size);
10    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
11
12    char processor_name[MPI_MAX_PROCESSOR_NAME];
13    int name_len;
14    MPI_Get_processor_name(processor_name, &name_len);
15
16    printf("Hello from processor %s, rank %d out of %d\n",
17          processor_name, world_rank, world_size);
```

```

17     processor_name, world_rank, world_size);
18
19     MPI_Finalize();
20     return 0;
21 }
```

Step 2: Complete execution workflow

```

1 # Connect to burst
2 ssh ab1234@gw.hpc.nyu.edu
3 ssh ab1234@greene.hpc.nyu.edu
4 ssh burst
5
6 # Get compute node with 8 CPUs
7 srun --account=csci_ua_0480_051-2025fa --partition=n2c48m24 --cpus-per-task=8 --time
   =04:00:00 --pty /bin/bash
8
9 # Navigate to scratch
10 cd /scratch/ab1234
11
12 # Launch Singularity (assuming you did the one-time setup)
13 singularity exec --overlay overlay-15GB-500K.ext3:ro ubuntu-20.04.3.sif /bin/bash
14
15 # Activate conda
16 source /ext3/miniconda3/bin/activate
17
18 # Create the program
19 vi hello_mpi.c
20 # (paste the code above)
21
22 # Compile
23 mpicc -o hello_mpi hello_mpi.c
24
25 # Run with 2 processes
26 mpirun -np 2 ./hello_mpi
27
28 # Expected output:
29 # Hello from processor b-9-77, rank 0 out of 8
30 # Hello from processor b-9-77, rank 1 out of 8
31 # Hello from processor b-9-77, rank 2 out of 8
32 # ... (8 lines total)
33
34 # Exit when done
35 exit # Exit Singularity
36 exit # Exit compute node
```

8.2 Example 2: CUDA Vector Addition

```

1 // vector_add.cu
2 #include <stdio.h>
3 #include <cuda_runtime.h>
4
5 __global__ void vectorAdd(float *a, float *b, float *c, int n) {
6     int i = blockDim.x * blockIdx.x + threadIdx.x;
7     if (i < n) {
8         c[i] = a[i] + b[i];
9     }
}
```

```

10 }
11
12 int main() {
13     int n = 1000;
14     size_t size = n * sizeof(float);
15
16     // Allocate host memory
17     float *h_a = (float*)malloc(size);
18     float *h_b = (float*)malloc(size);
19     float *h_c = (float*)malloc(size);
20
21     // Initialize vectors
22     for (int i = 0; i < n; i++) {
23         h_a[i] = i;
24         h_b[i] = i * 2;
25     }
26
27     // Allocate device memory
28     float *d_a, *d_b, *d_c;
29     cudaMalloc(&d_a, size);
30     cudaMalloc(&d_b, size);
31     cudaMalloc(&d_c, size);
32
33     // Copy to device
34     cudaMemcpy(d_a, h_a, size, cudaMemcpyHostToDevice);
35     cudaMemcpy(d_b, h_b, size, cudaMemcpyHostToDevice);
36
37     // Launch kernel
38     int threadsPerBlock = 256;
39     int blocksPerGrid = (n + threadsPerBlock - 1) / threadsPerBlock;
40     vectorAdd<<<blocksPerGrid, threadsPerBlock>>>(d_a, d_b, d_c, n);
41
42     // Copy result back
43     cudaMemcpy(h_c, d_c, size, cudaMemcpyDeviceToHost);
44
45     // Verify result
46     printf("First 5 results:\n");
47     for (int i = 0; i < 5; i++) {
48         printf("%f + %f = %f\n", h_a[i], h_b[i], h_c[i]);
49     }
50
51     // Cleanup
52     cudaFree(d_a);
53     cudaFree(d_b);
54     cudaFree(d_c);
55     free(h_a);
56     free(h_b);
57     free(h_c);
58
59     return 0;
60 }
```

To run:

```

1 # Get GPU node (NO Singularity needed for CUDA)
2 srun --account=csci_ua_0480_051-2025fa --partition=g2-standard-12 --gres=gpu:1 --time
   =02:00:00 --pty /bin/bash
3
4 # Load CUDA
```

```

5 module load cuda/11.8.0
6
7 # Compile and run
8 nvcc -o vector_add vector_add.cu
9 ./vector_add

```

8.3 Example 4: Complete Singularity Setup from Scratch

This shows the complete one-time setup process:

```

1 # 1. Connect to burst and get compute node
2 ssh ab1234@gw.hpc.nyu.edu
3 ssh ab1234@greene.hpc.nyu.edu
4 ssh burst
5 srun --account=csci_ua_0480_051-2025fa --partition=n2c48m24 --cpus-per-task=8 --time
   =02:00:00 --pty /bin/bash
6
7 # 2. Go to scratch directory
8 cd /scratch/ab1234
9
10 # 3. Download Singularity image
11 scp -rp greene-dtn:/scratch/work/public/singularity/ubuntu-20.04.3.sif .
12
13 # 4. Download and extract overlay
14 scp -rp greene-dtn:/scratch/work/public/overlay-fs-ext3/overlay-15GB-500K.ext3.gz .
15 gunzip overlay-15GB-500K.ext3.gz
16
17 # 5. Launch Singularity with read-write overlay
18 singularity exec --overlay overlay-15GB-500K.ext3:rw ubuntu-20.04.3.sif /bin/bash
19
20 # 6. Inside Singularity: Install Miniconda
21 wget https://repo.anaconda.com/miniconda/Miniconda3-latest-Linux-x86_64.sh
22 bash Miniconda3-latest-Linux-x86_64.sh -b -p /ext3/miniconda3
23
24 # 7. Activate conda
25 source /ext3/miniconda3/bin/activate
26
27 # 8. Install OpenMPI
28 conda install -c conda-forge openmpi -y
29
30 # 9. Test it works
31 echo '#include <mpi.h>
32 #include <stdio.h>
33 int main(int argc, char** argv) {
34     MPI_Init(&argc, &argv);
35     int rank;
36     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
37     printf("Hello from rank %d\n", rank);
38     MPI_Finalize();
39     return 0;
40 }' > test_mpi.c
41
42 mpicc -o test_mpi test_mpi.c
43 mpirun -np 4 ./test_mpi
44
45 # 10. If you see 4 lines of output, setup is complete!
46 # Exit and use :ro for future sessions
47 exit

```

9 Batch Job Scripts

9.1 MPI Batch Job

```
1 #!/bin/bash
2 #SBATCH --job-name=mpi_job
3 #SBATCH --account=csci_ua_0480_051-2025fa
4 #SBATCH --partition=n2c48m24
5 #SBATCH --nodes=1
6 #SBATCH --ntasks-per-node=32
7 #SBATCH --time=01:00:00
8 #SBATCH --output=%j_%x.out
9 #SBATCH --error=%j_%x.err
10 #SBATCH --requeue
11
12 # Navigate to scratch
13 cd /scratch/$USER
14
15 # Launch Singularity with conda and run program
16 singularity exec --overlay overlay-15GB-500K.ext3:ro \
17     ubuntu-20.04.3.sif /bin/bash << 'EOF'
18 source /ext3/miniconda3/bin/activate
19 mpirun -np 32 ./my_mpi_program
20 EOF
```

9.2 CUDA Batch Job

```
1 #!/bin/bash
2 #SBATCH --job-name=cuda_job
3 #SBATCH --account=csci_ua_0480_051-2025fa
4 #SBATCH --partition=g2-standard-12
5 #SBATCH --gres=gpu:1
6 #SBATCH --time=01:00:00
7 #SBATCH --output=%j_%x.out
8 #SBATCH --error=%j_%x.err
9 #SBATCH --requeue
10
11 module load cuda/11.8.0
12
13 echo "GPU Information:"
14 nvidia-smi
15 echo ""
16
17 echo "Running CUDA program..."
18 ./my_cuda_program
```

10 Troubleshooting

10.1 Common MPI Errors and Solutions

10.1.1 Error: libpmix.so.2: cannot open shared object file

Problem: You tried to run MPI programs directly without Singularity.

Solution: Always use Singularity with conda for MPI programs:

```
1 singularity exec --overlay overlay-15GB-500K.ext3:ro ubuntu-20.04.3.sif /bin/bash
2 source /ext3/miniconda3/bin/activate
3 mpirun -np 8 ./program
```

10.1.2 FATAL: while loading overlay images: failed to open overlay image

Problem: The overlay file might have a stale lock from a previous session that didn't exit cleanly.

Solution: Reboot Your Session:

```
1 mv overlay-15GB-500K.ext3 overlay-15GB-500K.ext3.backup
2 scp -rp greene-dtn:/scratch/work/public/overlay-fs-ext3/overlay-15GB-500K.ext3.gz .
3 gunzip overlay-15GB-500K.ext3.gz
4 singularity exec --overlay overlay-15GB-500K.ext3:rw ubuntu-20.04.3.sif /bin/bash
```

10.1.3 Error: Not enough slots available

Problem: You didn't request enough CPUs when launching your job.

Solution: Match CPU count to MPI processes:

```
1 # If running mpirun -np 8, need at least 8 CPUs:
2 srun --account=csci_ua_0480_051-2025fa --partition=n2c48m24 --cpus-per-task=8 --time
   =04:00:00 --pty /bin/bash
```

10.1.4 Error: File not found when copying Singularity image

Problem: Wrong file path.

Solution: Use the correct path:

```
1 # CORRECT path:
2 scp -rp greene-dtn:/scratch/work/public/singularity/ubuntu-20.04.3.sif .
3
4 # NOT this:
5 # scp -rp greene-dtn:/share/apps/images/ubuntu-20.04.3.sif .
```

10.1.5 Error: command not found: conda

Problem: Forgot to activate conda inside Singularity.

Solution: Always run after entering Singularity:

```
1 source /ext3/miniconda3/bin/activate
```

10.2 For CUDA Programs

10.2.1 Error: nvcc: command not found

Solution: Load CUDA module first:

```
1 module load cuda/11.8.0
```

All "module load" commands are for greene and not burst.

10.2.2 Error: No CUDA-capable device

Problem: Not on a GPU node.

Solution: Request GPU node:

```
1 srun --account=csci_ua_0480_051-2025fa --partition=g2-standard-12 --gres=gpu:1 --time
   =04:00:00 --pty /bin/bash
```

10.3 SSH Host Key Issues

If you see “Host key verification failed” when copying from greene-dtn:

```
1 # Remove old host key
2 ssh-keygen -R greene-dtn
3 ssh-keygen -R greene-dtn.hpc.nyu.edu
4
5 # Then retry your scp command
```

11 Debugging Tips

11.1 For MPI Programs

- **Start with 1 process:** `mpirun -np 1 ./program` to verify basic functionality
- **Add rank-specific output:** Help identify which process has issues
- **Use printf with fflush:** Ensures output appears immediately
- **Check MPI return codes:** Many MPI functions return error codes

Example debugging code:

```
1 printf("Rank %d: Before MPI_Send\n", rank);
2 fflush(stdout);
3 int result = MPI_Send(...);
4 if (result != MPI_SUCCESS) {
5     printf("Rank %d: MPI_Send failed with code %d\n", rank, result);
6 }
7 printf("Rank %d: After MPI_Send\n", rank);
8 fflush(stdout);
```

11.2 For CUDA Programs

- **Check every CUDA call:** Don’t skip error checking
- **Use cuda-memcheck:** Detects memory errors
- **Start with CPU version:** Verify algorithm before GPU
- **Use small test cases:** Easier to debug

Example error checking:

```
1 cudaError_t err = cudaMalloc(&d_array, size);
2 if (err != cudaSuccess) {
3     printf("CUDA Error: %s\n", cudaGetErrorString(err));
4     exit(1);
5 }
6
7 // After kernel launch
8 err = cudaGetLastError();
9 if (err != cudaSuccess) {
```

```

10     printf("Kernel launch error: %s\n", cudaGetErrorString(err));
11 }
```

11.3 General Debugging

- **Compile with debug symbols:** -g flag
- **Check you're in the right environment:** Inside Singularity for MPI, GPU node for CUDA
- **Verify file locations:** Use ls and pwd
- **Check resource allocation:** Use squeue -u \$USER

12 Time and Resource Management

12.1 GPU Hour Budget

With 300 GPU hours for the semester:

- **Development:** 15-30 min sessions for debugging
- **Testing:** 1 hour sessions for testing
- **Production:** 4+ hour sessions for final runs
- **Monitor usage:** Check regularly with sacct

Check your usage:

```

1 # View recent jobs
2 sacct -u $USER --format=JobID,JobName,Partition,Elapsed,State
3
4 # View all jobs since semester start
5 sacct -u $USER --format=JobID,JobName,Elapsed,State,AllocGRES -S 2025-01-01
```

12.2 Best Practices

- **Test interactively first:** Debug before submitting batch jobs
- **Use appropriate partitions:** CPU for MPI, GPU for CUDA
- **Request appropriate time:** Don't request 4 hours if you only need 30 minutes
- **Use --requeue:** Allows jobs to restart if nodes shut down
- **Save checkpoints:** For long-running jobs

13 Quick Reference

13.1 Essential Commands

```

1 # SSH to burst
2 ssh <netid>@gw.hpc.nyu.edu
3 ssh <netid>@greene.hpc.nyu.edu
4 ssh burst
5
6 # Interactive MPI job (with Singularity)
7 srun --account=csci_ua_0480_051-2025fa --partition=n2c48m24 --cpus-per-task=8 --time
   =04:00:00 --pty /bin/bash
8 cd /scratch/$USER
9 singularity exec --overlay overlay-15GB-500K.ext3:ro ubuntu-20.04.3.sif /bin/bash
10 source /ext3/miniconda3/bin/activate
11
```

```

12 # Interactive GPU job
13 srun --account=csci_ua_0480_051-2025fa --partition=g2-standard-12 --gres=gpu:1 --time
   =04:00:00 --pty /bin/bash
14 module load cuda/11.8.0
15
16 # Submit batch job
17 sbatch job_script.sh
18
19 # Check job status
20 squeue -u $USER
21
22 # Cancel job
23 scancel <job_id>
24
25 # View job history
26 sacct -u $USER

```

13.2 MPI Workflow Summary

```

1 # 1. Get compute node with CPUs
2 srun --account=csci_ua_0480_051-2025fa --partition=n2c48m24 --cpus-per-task=8 --time
   =04:00:00 --pty /bin/bash
3
4 # 2. Navigate to scratch
5 cd /scratch/<netid>
6
7 # 3. Launch Singularity
8 singularity exec --overlay overlay-15GB-500K.ext3:ro ubuntu-20.04.3.sif /bin/bash
9
10 # 4. Activate conda
11 source /ext3/miniconda3/bin/activate
12
13 # 5. Compile and run
14 mpicc -o program program.c
15 mpirun -np 2 ./program

```

13.3 CUDA Workflow Summary

```

1 # 1. Get GPU node
2 srun --account=csci_ua_0480_051-2025fa --partition=g2-standard-12 --gres=gpu:1 --time
   =04:00:00 --pty /bin/bash
3
4 # 2. Load CUDA
5 module load cuda/11.8.0
6
7 # 3. Compile and run
8 nvcc -o program program.cu
9 ./program

```

14 Summary

14.1 Key Takeaways

- **Singularity is REQUIRED for MPI:** Burst nodes have broken MPI, use Singularity + conda

- **Singularity NOT needed for CUDA:** CUDA works directly with module load
- **Specify CPUs for MPI:** Use `--cpus-per-task=N` matching your MPI process count
- **Use correct file paths:** `/scratch/work/public/singularity/...`
- **Activate conda in Singularity:** `source /ext3/miniconda3/bin/activate`
- **Work in /scratch:** Better performance than `/home`
- **Open OnDemand is easiest:** <https://ood-burst-001.hpc.nyu.edu/>
- **Use `--requeue`:** In all batch scripts for spot instances
- **Monitor GPU hours:** 300 hours total for semester
- **Test interactively first:** Debug before batch submission

14.2 Recommended Workflow

1. Develop code locally or using VS Code Remote-SSH
2. Upload to HPC via Open OnDemand or SCP
3. For MPI: Set up Singularity once, then reuse
4. Test interactively with small inputs
5. Debug and verify correctness
6. Create batch script for production runs
7. Submit batch job with `--requeue`
8. Monitor with `squeue` and download results

14.3 Important Links

- **Open OnDemand:** <https://ood-burst-001.hpc.nyu.edu/>
- **NYU HPC Docs:** <https://sites.google.com/nyu.edu/nyu-hpc/>
- **Singularity Guide:** <https://sites.google.com/nyu.edu/nyu-hpc/hpc-systems/greene/software/singularity-with-miniconda>
- **HPC Support:** `hpc@nyu.edu`

Good Luck!

You now have everything you need to successfully use HPC Cloud Bursting! Remember:

- MPI requires Singularity (one-time setup, then easy to use)
- CUDA works directly without Singularity
- Always specify CPUs when running MPI programs
- Start early and test often
- Ask questions when stuck - check Brightspace or email `hpc@nyu.edu`
- Have fun learning parallel computing!